

Saynam Sharma

Data Engineer

✉ saynam1101@gmail.com ☎ +91 9419271101  LinkedIn  GitHub  Portfolio

PROFILE

Data Engineer with 3.5+ years of experience building and owning scalable batch and real-time data pipelines on Databricks, AWS, and Snowflake. Delivered 45–60% runtime improvements through Spark optimization and designed resilient ETL architectures handling millions of records per run. Specialized in incremental processing, Delta Lake patterns, schema evolution, and pipeline observability for production-grade systems.

PROFESSIONAL EXPERIENCE

KPI Partners, *Data Engineer II*

09/2022 – present | Remote- Pune, India

Oracle Fusion → Snowflake ETL Platform (AWS, Airflow, dbt)

- Architected an end-to-end AWS-based ETL platform ingesting Oracle Fusion BI report data into Snowflake via MWA (Airflow), S3, Python, and SQL — enabling analytics-ready datasets for enterprise reporting across multiple business domains.
- Designed incremental and idempotent load strategies transforming raw Fusion data into dimensional and fact models using dbt on Snowflake, ensuring safe reruns and consistent downstream data states.
- Built dbt transformation layers with embedded data quality checks, audit columns, and schema drift handling, reducing downstream reporting failures and improving data reliability.
- Engineered backfill-safe execution mechanisms supporting full historical reloads and partial reprocessing without impacting live production data.
- Enhanced pipeline observability through structured logging, audit tables, and Snowflake query history monitoring — cutting mean time to resolution for pipeline failures.

Databricks & PySpark Lakehouse Pipelines (Historical + Incremental Data)

- Built large-scale batch pipelines on Databricks using PySpark and Spark SQL to process and merge historical DBCS datasets with incoming Fusion source data across multi-million-row fact tables.
- Implemented client-defined merge and transformation logic using Delta Lake patterns to reliably handle late-arriving data, malformed records, and schema changes without data corruption.
- Developed complex PySpark transformations leveraging window functions, conditional ordering, and deduplication logic — ensuring correctness and scalability across large fact tables.
- Optimized Spark workloads by redesigning partitioning strategies, eliminating unnecessary shuffles, and tuning Databricks cluster configurations — achieving 45–60% reduction in pipeline runtime.
- Stabilized flaky production pipelines by introducing structured error handling, rerun-safe mechanisms, and audit columns — significantly reducing on-call incidents and manual reprocessing effort.

Clinical Trials Data Platform (SSIS, SSMS, Databricks)

- Developed and maintained legacy SSIS pipelines to ingest and process clinical trial data, managing trial IDs and supporting end-to-end data flow from source systems to frontend-facing applications.
- Built and optimized SQL Server stored procedures to enforce business logic, manage trial data transformations, and ensure reliable data delivery across the clinical data platform.
- Designed structured data workflows in SSMS to handle trial lifecycle stages, ensuring data integrity and consistency across downstream systems.
- Currently leading migration of legacy SSIS workflows to Databricks pipelines, modernizing the data platform from on-premise SQL Server architecture to a cloud-native lakehouse stack.

EDUCATION

Sir M Visvesvaraya Institute of Technology,
B.E in Electronics and Communication

08/2022 | Bangalore, India

SKILLS

Python | Scala | SQL | AWS | Apache Spark | AirFlow | Snowflake | Git | Databricks |
Hadoop | Docker | Pandas | Azure

CERTIFICATES

Databricks Certified Data Engineer Associate | Databricks Certified Data Engineer Professional |
Python Programming | Apache Spark Lake House Fundamentals